

```
In [15]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
In [18]: # Dataset Load
customers = pd.read_csv('C:/Users/avina/Downloads/Zeotap_Assignment/Customers.csv')
products = pd.read_csv('C:/Users/avina/Downloads/Zeotap_Assignment/Products.csv')
transactions = pd.read_csv('C:/Users/avina/Downloads/Zeotap_Assignment/Transactions.csv')
```

```
In [19]: #Few rows
print("Customers:")
print(customers.head(), '\n')
print("Products:")
print(products.head(), '\n')
print("Transactions:")
print(transactions.head(), '\n')
```

Customers:

	CustomerID	CustomerName	Region	SignupDate
0	C0001	Lawrence Carroll	South America	2022-07-10
1	C0002	Elizabeth Lutz	Asia	2022-02-13
2	C0003	Michael Rivera	South America	2024-03-07
3	C0004	Kathleen Rodriguez	South America	2022-10-09
4	C0005	Laura Weber	Asia	2022-08-15

Products:

	ProductID	ProductName	Category	Price
0	P001	ActiveWear Biography	Books	169.30
1	P002	ActiveWear Smartwatch	Electronics	346.30
2	P003	ComfortLiving Biography	Books	44.12
3	P004	BookWorld Rug	Home Decor	95.69
4	P005	TechPro T-Shirt	Clothing	429.31

Transactions:

	TransactionID	CustomerID	ProductID	TransactionDate	Quantity	\
0	T00001	C0199	P067	2024-08-25 12:38:23	1	
1	T00112	C0146	P067	2024-05-27 22:23:54	1	
2	T00166	C0127	P067	2024-04-25 07:38:55	1	
3	T00272	C0087	P067	2024-03-26 22:55:37	2	
4	T00363	C0070	P067	2024-03-21 15:10:10	3	

	TotalValue	Price
0	300.68	300.68
1	300.68	300.68
2	300.68	300.68
3	601.36	300.68
4	902.04	300.68

```
In [20]: #Data Inspection
print(customers.info())
print(products.info())
print(transactions.info())
```

#Check null values

```
print(customers.isnull().sum())
```

```
print(products.isnull().sum())
```

```
print(transactions.isnull().sum())
```

#Summary

```
print(transactions.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   CustomerID      200 non-null   object
1   CustomerName    200 non-null   object
2   Region          200 non-null   object
3   SignupDate      200 non-null   object
```

dtypes: object(4)

memory usage: 6.4+ KB

None

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   ProductID       100 non-null   object
1   ProductName     100 non-null   object
2   Category        100 non-null   object
3   Price           100 non-null   float64
```

dtypes: float64(1), object(3)

memory usage: 3.3+ KB

None

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   TransactionID    1000 non-null   object
1   CustomerID       1000 non-null   object
2   ProductID        1000 non-null   object
3   TransactionDate   1000 non-null   object
4   Quantity         1000 non-null   int64
5   TotalValue       1000 non-null   float64
6   Price            1000 non-null   float64
```

dtypes: float64(2), int64(1), object(4)

memory usage: 54.8+ KB

None

CustomerID 0

CustomerName 0

Region 0

SignupDate 0

dtype: int64

ProductID 0

ProductName 0

Category 0

Price 0

dtype: int64

TransactionID 0

CustomerID 0

ProductID 0

TransactionDate 0

Quantity 0


TotalValue 0

Price 0


```
dtype: int64
```

	Quantity	TotalValue	Price
count	1000.000000	1000.000000	1000.000000
mean	2.537000	689.995560	272.55407
std	1.117981	493.144478	140.73639
min	1.000000	16.080000	16.080000
25%	2.000000	295.295000	147.95000
50%	3.000000	588.880000	299.93000
75%	4.000000	1011.660000	404.40000
max	4.000000	1991.040000	497.76000

```
In [60]: # Visualizations_1
# Signups (By region)
sns.countplot(data=customers, x='Region', hue='Region', palette='viridis', legend=F
plt.title('Number of Customers by Region')
plt.show()
```



```
In [22]: #Visualization_2
#Product prices distribution
sns.histplot(data=products, x='Price', bins=20, kde=True)
plt.title('Distribution of Product Prices')
plt.xlabel('Price (USD)')
plt.show()
```




```
In [26]: #Visualization_3
#Transactions v/s time

transactions['TransactionDate'] = pd.to_datetime(transactions['TransactionDate'])

transactions['Month'] = transactions['TransactionDate'].dt.to_period('M')

monthly_sales = transactions.groupby('Month')['TotalValue'].sum()

# Plot the monthly sales trend
plt.figure(figsize=(10, 6))
plt.plot(monthly_sales.index.astype(str), monthly_sales.values, marker='o')
plt.title('Monthly Sales Trend')
plt.xlabel('Month')
plt.ylabel('Total Sales (USD)')
plt.xticks(rotation=45)
plt.grid()
plt.tight_layout()
plt.show()
```



```
In [29]: # Prepare data by merging datasets

#MERGING DATASETS

data = transactions.merge(customers, on='CustomerID').merge(products, on='ProductID')

print(data.columns)

#Profile creation of each customer
customer_profiles = data.groupby('CustomerID').agg({
    # Average product price purchased
    'Price_y': 'mean',
    # Total quantity purchased
    'Quantity': 'sum',
    # Total spend
    'TotalValue': 'sum'
}).reset_index()

print(customer_profiles.head())
```

```
Index(['TransactionID', 'CustomerID', 'ProductID', 'TransactionDate',
      'Quantity', 'TotalValue', 'Price_x', 'Month', 'CustomerName', 'Region',
      'SignupDate', 'ProductName', 'Category', 'Price_y'],
      dtype='object')
   CustomerID  Price_y  Quantity  TotalValue
0      C0001  278.334000         12      3354.52
1      C0002  208.920000         10      1862.74
2      C0003  195.707500         14      2725.38
3      C0004  240.636250         23      5354.88
4      C0005  291.603333          7      2034.24
```

```
In [32]: #Similarity Calculation (cosine similarity)

import os
```

```

import csv
from sklearn.metrics.pairwise import cosine_similarity
import pandas as pd

# Assumed -> Customer_profiles DataFrame already created

# Prepare the input features: we can use average price, total spend, total quantity
X = customer_profiles[['Price_y', 'Quantity', 'TotalValue']].values

# Compute cosine similarity
similarity_matrix = cosine_similarity(X)
results = {}

# Iterate over customers to find their Lookalikes
for i, customer_id in enumerate(customer_profiles['CustomerID']):

    similarities = list(enumerate(similarity_matrix[i]))
    similarities = sorted(similarities, key=lambda x: x[1], reverse=True)[1:4]

    results[customer_id] = [(customer_profiles.iloc[idx, 0], round(score, 2)) for i, (idx, score) in similarities]

os.makedirs('outputs', exist_ok=True)

with open('outputs/Lookalike.csv', 'w', newline='') as f:
    writer = csv.writer(f)
    writer.writerow(['CustomerID', 'Lookalikes'])

    for customer_id, lookalikes in results.items():
        # Format the list of lookalikes for writing
        lookalike_str = ', '.join([f'{lookalike[0]} ({lookalike[1]})' for lookalike in lookalikes])
        writer.writerow([customer_id, lookalike_str])

print("Lookalike model results written to 'outputs/Lookalike.csv'")

```

Lookalike model results written to 'outputs/Lookalike.csv'

In [39]: `print(data.columns)`

```

Index(['TransactionID', 'CustomerID', 'ProductID', 'TransactionDate',
      'Quantity', 'TotalValue', 'Price_x', 'Month', 'CustomerName', 'Region',
      'SignupDate', 'ProductName', 'Category', 'Price_y'],
      dtype='object')

```

In [45]: `print(data.head())`

	TransactionID	CustomerID	ProductID	TransactionDate	Quantity	\
0	T00001	C0199	P067	2024-08-25 12:38:23	1	
1	T00112	C0146	P067	2024-05-27 22:23:54	1	
2	T00166	C0127	P067	2024-04-25 07:38:55	1	
3	T00272	C0087	P067	2024-03-26 22:55:37	2	
4	T00363	C0070	P067	2024-03-21 15:10:10	3	

	TotalValue	Price_x	Month	CustomerName	Region	SignupDate	\
0	300.68	300.68	2024-08	Andrea Jenkins	Europe	2022-12-03	
1	300.68	300.68	2024-05	Brittany Harvey	Asia	2024-09-04	
2	300.68	300.68	2024-04	Kathryn Stevens	Europe	2024-04-04	
3	601.36	300.68	2024-03	Travis Campbell	South America	2024-04-11	
4	902.04	300.68	2024-03	Timothy Perez	Europe	2022-03-15	


	ProductName	Category	Price_y
0	ComfortLiving Bluetooth Speaker	Electronics	300.68
1	ComfortLiving Bluetooth Speaker	Electronics	300.68
2	ComfortLiving Bluetooth Speaker	Electronics	300.68
3	ComfortLiving Bluetooth Speaker	Electronics	300.68
4	ComfortLiving Bluetooth Speaker	Electronics	300.68

```
In [46]: from sklearn.cluster import KMeans
from sklearn.metrics import davies_bouldin_score

# Davies-Bouldin Index
# Optimal number of clusters
scores = []
# Check for k from 2 to 10
for k in range(2, 11):
    kmeans = KMeans(n_clusters=k, random_state=42)
    clusters = kmeans.fit_predict(features)
    score = davies_bouldin_score(features, clusters)
    scores.append(score)

# Plot Index for each k
plt.plot(range(2, 11), scores, marker='o')
plt.title('Davies-Bouldin Index vs Number of Clusters')
plt.xlabel('Number of Clusters')
plt.ylabel('Davies-Bouldin Index')
plt.show()

# Lowest Davies-Bouldin index with one optimal k
optimal_k = scores.index(min(scores)) + 2
print(f"Optimal number of clusters: {optimal_k}")
```

Optimal number of clusters: 2

```
In [47]: # Apply K-Means with optimal number of clusters
kmeans = KMeans(n_clusters=optimal_k, random_state=42)
customer_profiles['Cluster'] = kmeans.fit_predict(features)

print(customer_profiles.head())
```

	CustomerID	Price_y	Quantity	TotalValue	Cluster
0	C0001	278.334000	12	3354.52	1
1	C0002	208.920000	10	1862.74	1
2	C0003	195.707500	14	2725.38	1
3	C0004	240.636250	23	5354.88	0
4	C0005	291.603333	7	2034.24	1

```
In [58]: import matplotlib.pyplot as plt
from sklearn.decomposition import PCA

# 2D for visualization
pca = PCA(n_components=2)
reduced_features = pca.fit_transform(features)

# DataFrame to hold the 2D PCA components and cluster labels
customer_profiles['PCA1'] = reduced_features[:, 0]
customer_profiles['PCA2'] = reduced_features[:, 1]

# Get unique clusters
unique_clusters = customer_profiles['Cluster'].nunique()
```

```

# Define marker styles based on the number of clusters
markers = ['o', 's', 'D', 'X', 'P', '^', 'v', '<', '>'][unique_clusters]

# Plot the clusters
plt.figure(figsize=(9,6))


# Use Seaborn's scatterplot with dynamic markers
sns.scatterplot(data=customer_profiles, x='PCA1', y='PCA2', hue='Cluster', palette=
                style='Cluster', markers=markers, s=100, edgecolor='black', legend=

# Adding titles and Labels for clarity
plt.title('Customer Segmentation with K-Means Clustering', fontsize=16)
plt.xlabel('PCA Component 1', fontsize=12)
plt.ylabel('PCA Component 2', fontsize=12)

# Display the Legend
plt.legend(title='Cluster', bbox_to_anchor=(1.05, 1), loc='upper left')

# Show the plot with a clean layout
plt.tight_layout()
plt.show()

```



```

In [49]: # Calculate final Davies-Bouldin Index
db_index = davies_bouldin_score(features, customer_profiles['Cluster'])
print(f"Davies-Bouldin Index for the final clustering: {db_index}")

```

Davies-Bouldin Index for the final clustering: 0.8502362410640188

In []: